

CS118 Discussion 1B

Week 2

Zhiyi Zhang (zhiyi@cs.ucla.edu)

Contents

- Socket programming (Cont'd)
- Application Layer
- Review HTTP
- Review DNS

Socket Programming (cont'd)

Socket Programming

Continue our programming.

Plan:

- First review the server code
- Then work on the client code

Next two pages

- Page 5: TCP server sample code
- Page 6: TCP client sample code

```

// *** Author: Zhiyi Zhang for CS118, Time: 04/10/2020, TCP Server Sample Code ***
#include <stdio.h>
#include <errno.h>
#include <string.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/socket.h>
#include <netdb.h>
#include <netinet/in.h>
#include <arpa/inet.h>

int main()
{
    // *** Initialize socket for listening ***
    int sockfd;
    if ((sockfd = socket(PF_INET, SOCK_STREAM, 0)) == -1) {
        perror("socket");
        exit(1);
    }

    // *** Initialize local listening socket address ***
    struct sockaddr_in my_addr;
    memset(&my_addr, 0, sizeof(my_addr));
    my_addr.sin_family = AF_INET;
    my_addr.sin_port = htons(5678);
    my_addr.sin_addr.s_addr = htonl(INADDR_ANY); // INADDR_ANY allows to connect to any one of the host's IP address

    // *** Socket Bind ***
    if (bind(sockfd, (struct sockaddr *)&my_addr, sizeof(struct sockaddr)) == -1) {
        perror("bind");
        exit(1);
    }

    // *** Socket Listen ***
    if (listen(sockfd, 10) == -1) {
        perror("listen");
        exit(1);
    }
}

```

```

// *** Author: Zhiyi Zhang for CS118, Time: 04/10/2020, TCP Client Sample Code ***
#include <stdio.h>
#include <string.h>
#include <sys/socket.h>
#include <netdb.h>
#include <netinet/in.h>

int main()
{
    // *** Initialize Socket ***
    int sockfd;
    if ((sockfd = socket(PF_INET, SOCK_STREAM, 0)) == -1) {
        perror("Cannot create socket");
        exit(1);
    }

    // *** Initialize the server socket address ***
    struct sockaddr_in server_addr; // server socket address struct
    server_addr.sin_family = AF_INET; // protocol family
    server_addr.sin_port = htons(5678); // port number
    struct hostent *host_name = gethostbyname("localhost"); // get IP from host name
    server_addr.sin_addr = *((struct in_addr *)host_name->h_addr); // set IP address
    memset(server_addr.sin_zero, '\0', sizeof server_addr.sin_zero); // make the rest bytes zero

    // *** Connect to the server ***
    if (connect(sockfd, (struct sockaddr *)&server_addr, sizeof(struct sockaddr)) == -1) {
        perror("Cannot connect");
        exit(1);
    }

    // *** Socket Read & Write ***
    int sin_size, recvline_size;
    char sendline[1024], recvline[1024];
    while (fgets(sendline, 1024, stdin) != NULL) {
        write(sockfd, sendline, strlen(sendline));
        if (memcmp(sendline, "bye", strlen("bye")) == 0) {
            printf("Will close the connection\n");
            close(sockfd);
        }
    }
}

```

Application Layer




Fat layer with so many different applications

- Web: HTTP
- Email: SMTP, POP
- DNS
- P2P Applications
- Video streaming
- ... so many others

They are all enabled by a thin transport layer (TCP, UDP) and an even thinner layer (IP)

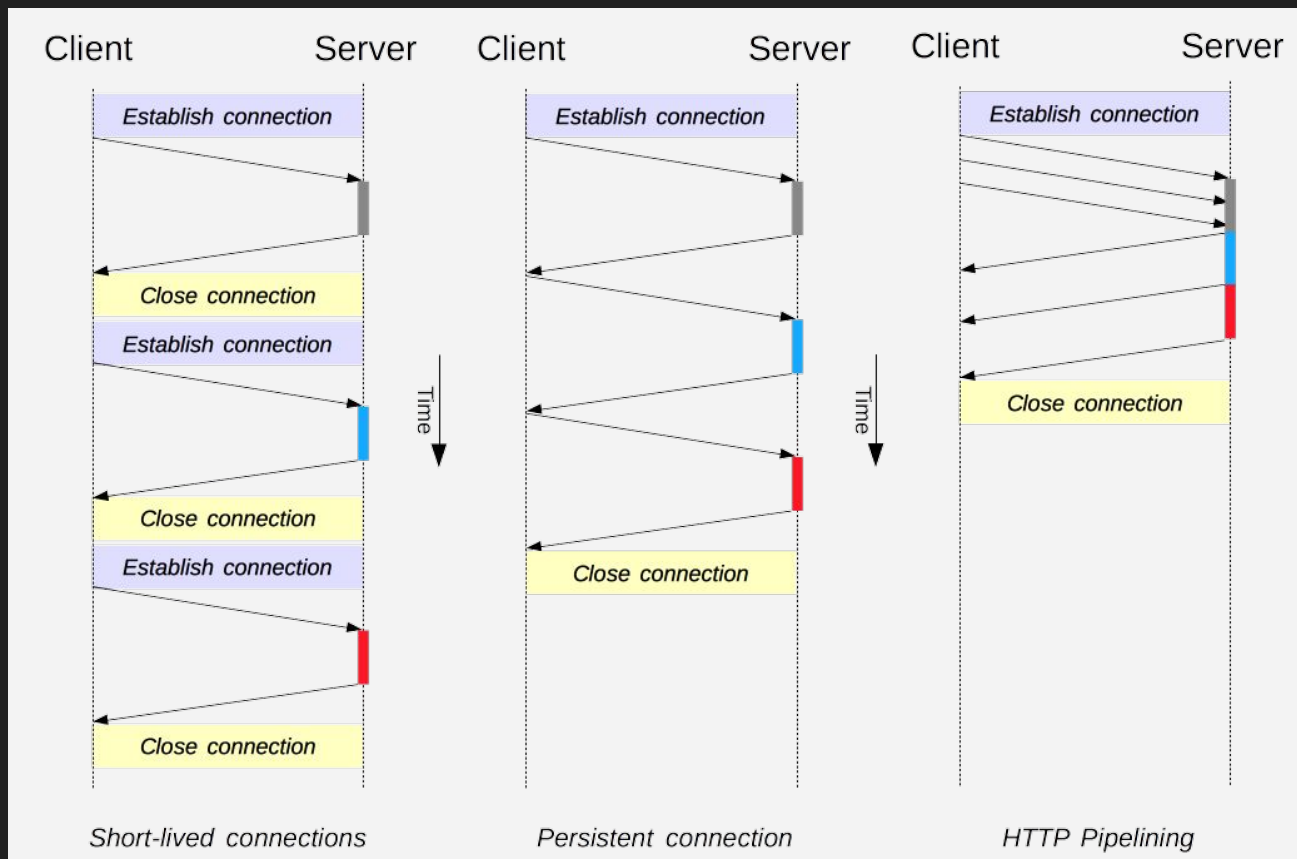
Models

- As we mentioned in previous discussion
 - Client-server: Web, DNS
 - P2P: BitTorrent, BitCoin, Onion Routing
 - Hybrid of two? Skype, BitTorrent + BT Site



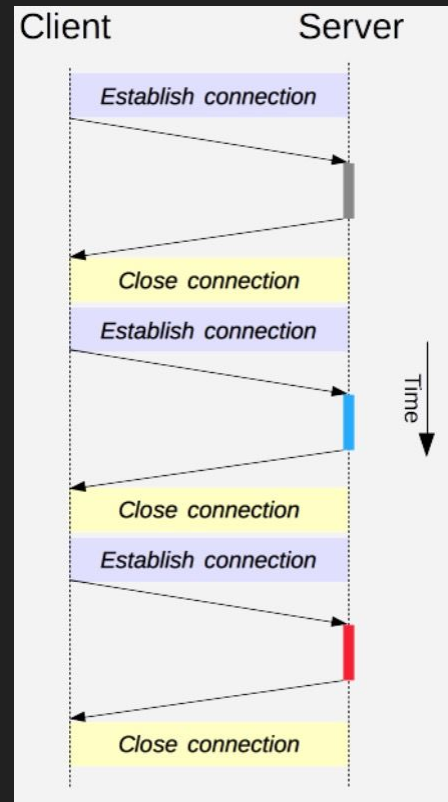
Review HTTP

Non-persistent v.s. Persistent v.s. Pipelining




Response Time: Non-persistent HTTP

- To fetch each object, HTTP requires
 - 1 RTT to set up the TCP connection
 - 1 RTT to get the data back
 - Data transmission time
- Therefore: one html with N objects
 - 2 RTTs + N*2 RTTs + Transmission Time
- Example question
 - One HTTP html page with 2 objects? (ignoring tx time)
 - HTTP page: 2 RTTs
 - 2 objects: 2*2 RTTs = 4 RTTs
 - In total: 6 RTTs




Response Time: Persistent HTTP without Pipeline

- Use the same TCP connection for other objects
- To fetch all objects, HTTP requires
 - 1 RTT to set up the TCP connection
 - 1 RTT to get the HTML file back
 - N RTT to get N objects back, e.g., CSS, JS, pictures
 - Data transmission time
- Therefore
 - $(2 + N)$ RTTs + Transmission Time
- Example question
 - One HTTP page with 2 objects on the same server? (ignoring tx time)
 - HTTP page: 2 RTTs
 - 2 objects: 2 RTTs = 2 RTTs
 - In total: 4 RTTs



Response Time: Persistent HTTP with Pipeline

- Use the same TCP connection for other objects
- To fetch each object, HTTP requires
 - 1 RTT to set up the TCP connection
 - 1 RTT to get the HTML file back
 - 1 RTT to get N objects back, e.g., CSS, JS, pictures
 - Data transmission time
- Therefore
 - 3 RTTs + Transmission Time
- Example question
 - One HTTP page with 2 objects on the same server? (ignoring tx time)
 - HTTP page: 2 RTTs
 - 2 objects: 1 RTTs
 - In total: 3 RTTs



Cookie makes HTTP stateful

- Components of Cookie

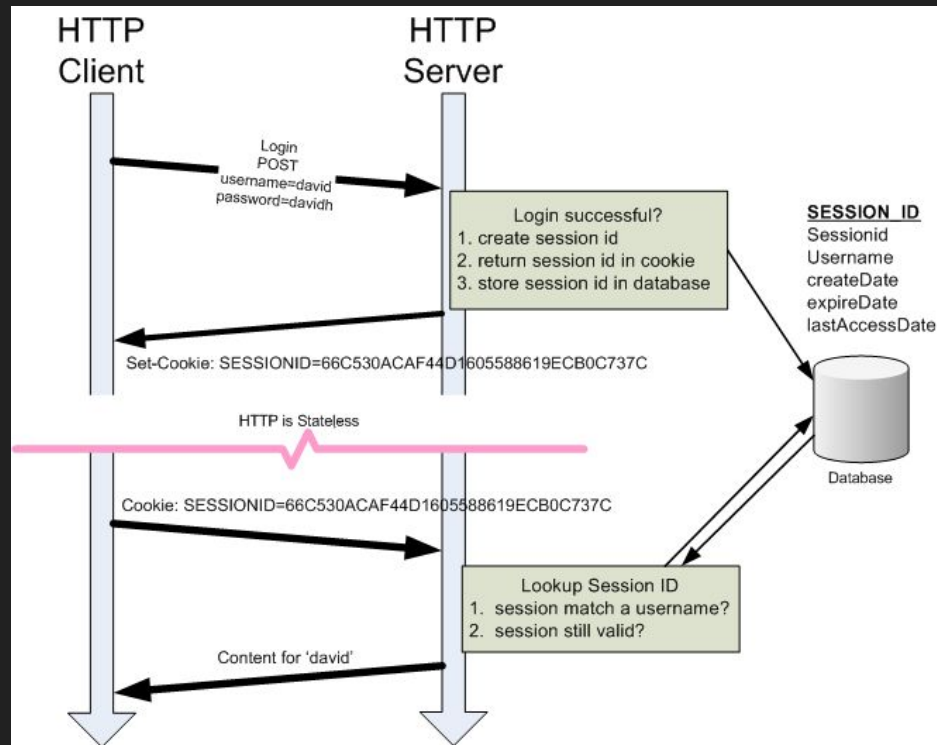
- Cookie header line in response: server initializes the cookie for the client
- Cookie header line in request: client specifies its cookie
- Cookie file on the browser
- Cookie stored in server database

- Main Purposes of Cookie

- Session management (e.g., your shopping cart keeps the same when you navigate the website, you are kept login for some time)
- Personalization (e.g., the website remembers your preferences)
- Tracking (e.g., advertisements with third-party cookie)

Cookie: A Typical Workflow

Cookie is kept by your browser



Where is our current Internet? (FYI)

- HTTP/1.1
- HTTP/2
- QUIC \leq HTTP/3

First need to understand the limitations of 1.1 (FYI)

- High page load latency
 - Head of line blocking
 - The response must be returned in full, in the order of request
 - Large http header
 - Responses are replied only after requesting them first

HTTP/2 (FYI)

- New features
 - Multiplexing over a single TCP connection
 - Responses can come back in arbitrary order
 - HTTP header compression
 - Server push
 - Server can push objects to the client without requests (pre-loading)
- Limitations
 - HTTPS = HTTP over TCP + TLS: When used with TLS, the handshake time is long (TCP handshake + TLS handshake)
 - Cannot work with mobility -- when client changes its IP address, TCP connection is broken

QUIC (Quick UDP Interaction Connections) (FYI)

- Over UDP
- QUIC realizes their own reliability, in-order packet delivery, congestion control plus
 - Combined TLS and QUIC handshake: 0 RTT
 - When QUIC finishes its handshake, so does TLS
 - Connection ID to identity connection instead of <src socket, dest socket>
 - When IP changes, your connection ID won't
- QUIC is in user space (kernel space)
 - Allow easy version update (evolvability)
 - Plug and play congestion control (modularity)

Review DNS

DNS components

- End hosts
 - Who wants to know the IP address of `www.ucla.edu`
- Local DNS server / DNS Resolver
 - Who helps remember existing DNS responses for the local network
- Authoritative Name servers
 - Who has the knowledge of Domain \Leftrightarrow IP mappings of its own zone


Types of Query

Transport layer used:

- Mainly on UDP

Iterative v.s. Recursive query

- Iterative query
 - Step 2-7
- Recursive query
 - Step 1 and 8



How does it scale for global use?

- Hierarchical structure
- Each authoritative name server cares their own business

dig command and nslookup command

See the DNS query and answer of a domain name.