

CS118 Discussion 1A, Week 1

Qianru Li

Friday / 10-11:50am

Outline

- Logistics
- Intro to network programming

TA

- Qianru Li, Ph.D. student in computer science
- Office hours: Monday 12:30 – 2:30 p.m.
- Email: qianruli@ucla.edu
 - If you did not get reply in 24 hrs, please send it again.
 - Please use [CS118] in subject to make your email stand out
 - Please include your name, UID in email.

Logistics

- Submit your signed Academic Integrity Agreement
- Grading breakdown
 - Homework 20%
 - Projects 20%
 - Up to 3% bonus for extra work
 - 4 quizzes: 60%

Logistics: Homework

- Online submission to Gradescope only (course entry code: **932KV3**). DEMO
 - Fill in your UCLA ID so that we can submit your score to CCLE
- Submission guidelines:
 - 1. **Hard deadline** on submission, so submit early! You can **resubmit** multiple times before the deadline, but the system will not accept submissions after the deadline.
 - 2. Each homework problem will have a dedicated **answering box** immediately below. Do **NOT** write your answers outside the box. Any answer outside the dedicated area may not get graded.
 - 3. You are encouraged to work out the problem on the PDF file directly **without altering the page layout in any way.**
 - 4. If you prefer handwriting or have to draw diagrams, you may scan the paper copy (e.g. using smartphone app like CamScanner) or take a picture, convert it to a PDF file and then upload. It is **your** responsibility to upload a high-quality copy in black and white. Inaccessible answers will get low scores.


Logistics: Project

- Two projects (in C/C++):
 - A simple web server — get familiar with network programming;
 - Reliable data transfer — implement a simple user-level TCP-like transport protocol
- Individual project.
- Test environment:
 - Ubuntu virtual machine

What are we learning in this course

- Part 1: Introduction (2 lectures, text: Chapter 1)
- Part 2: Application Layer (2 lectures, text: Ch.2)
 - *Quiz 1 to cover Parts 1 &2
- Part 3: Transport Layer (4.5 lectures, text Ch. 3)
 - *Project 1: due April 24, Friday
 - *Quiz 2 to cover Part 3
- Part 4: Network Layer (4 lectures, text: Ch. 4 and 5)
 - *Quiz 3 to cover Part 4
- Part 5: Link Layer, LANs (3.5 lectures, text: Ch. 6)
- Part 6: Wireless and Mobile Networks (1.5 lectures, text: Ch. 7)
- Part 8: Network Security (0.5 lecture: Ch. 8)
 - *Project 2: due June 5, Friday
 - *Quiz 4 to cover Parts 5, 6 & 8 (using final exam slot).

Network Programming



Network programming

- **What is the model for network programming?**
- At which layer do we program?
- Which APIs and usage?

Client-server model

- Asymmetric communication
 - Client — requests data:
 - Initiates communication
 - Waits for server's response
 - Server (Daemon) — responds data requests:
 - Discoverable by clients (e.g. IP address + port)
 - Waits for clients connection
 - Processes requests, sends replies

Client-server model

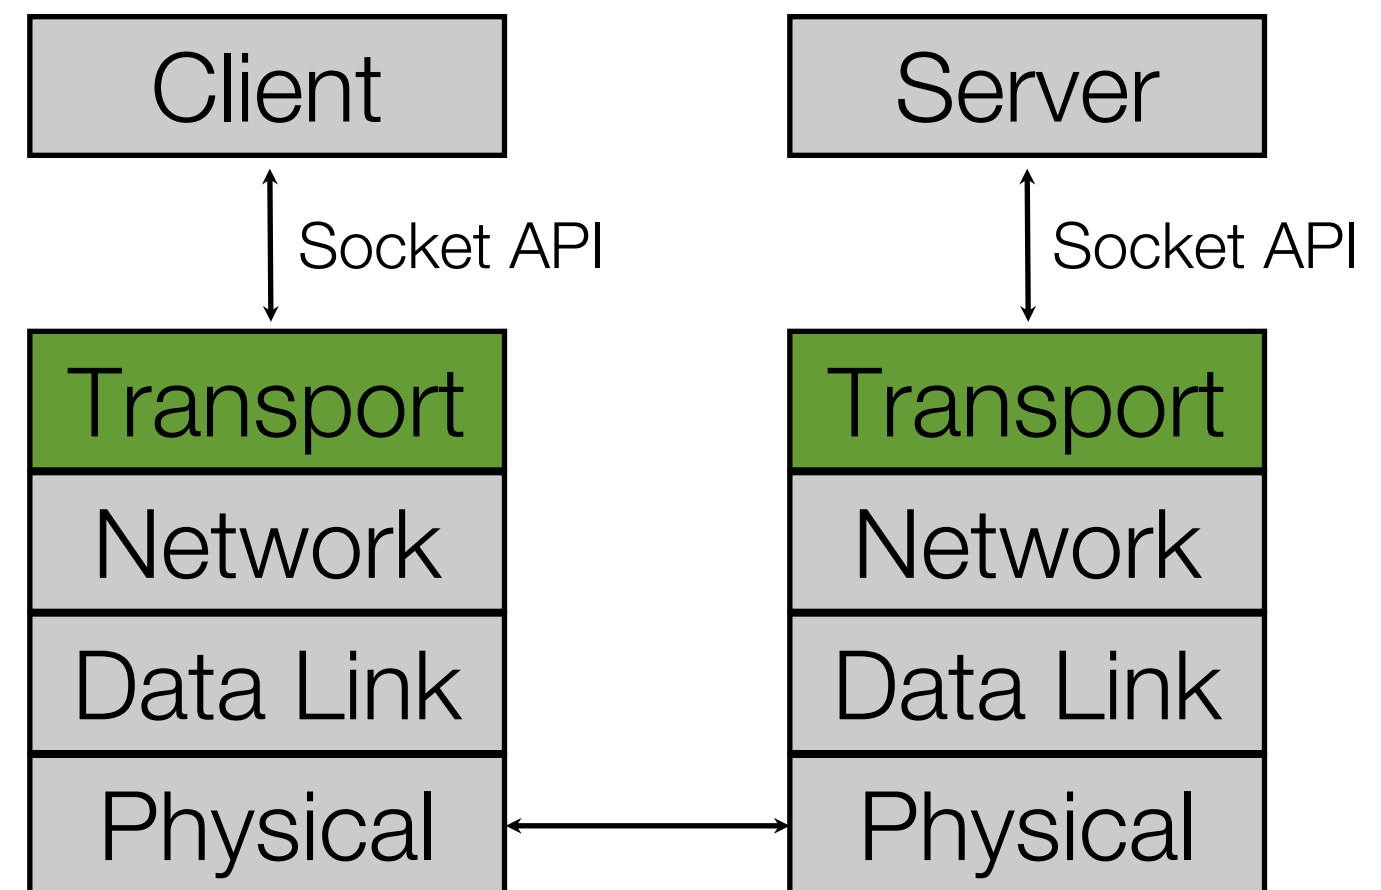
- Client and server are not disjoint
 - A client can be a server of another client
 - A server can be a client of another server
 - Example?
- Server's service model
 - Concurrent: server **processes** multiple clients' requests simultaneously
 - Sequential: server processes clients' requests one by one
 - Hybrid: server maintains multiple connections, but responses sequentially

Network programming

- What is the model for network programming?
- **At which layer do we program?**
- Which APIs and usage?

Which layer are we at?

- “Clients” and “servers” are programs at application layer
- Transport layer is responsible for providing communication services for application layer
- Basic transport layer protocols:
 - TCP
 - UDP



TCP: Transmission Control Protocol

- A connection is set up between client and server
- Reliable data transfer
 - Guarantee deliveries of all data
 - No duplicate data would be delivered to application
- Ordered data transfer
 - If Alice sends data D1 followed by D2 to Bob, Bob will also receive D1 before D2
- Data transmission: full-duplex byte stream (in two directions simultaneously)
- Regulated data flow: flow control and congestion control

UDP: User Data Protocol


- Basic data transmission service
 - Unit of data transfer: datagram (in variable length)
- No reliability guarantee
- No ordered delivery guarantee
- No flow control / congestion control

Network programming

- What is the model for network programming?
- At which layer do we program?
- **Which APIs and usage?**

Our secret weapon: socket programming APIs

- From Wikipedia: “A network socket is an endpoint of an inter-process communication flow across a computer network”
- A socket is a tuple of `<ip_addr:port>`
- Socket programming APIs help build the communication tunnel between applications and transport/network service
- We use TCP socket in this project



TCP socket: basic steps

- Create service
- Establish a TCP connection
- Send and receive data
- Close the TCP connection

TCP socket: service setup

TCP Client

TCP Server

TCP socket: service setup

TCP Client


TCP Server

socket()

TCP socket: service setup

TCP Client


TCP Server



TCP socket: service setup

TCP Client


TCP Server



TCP socket: service setup

TCP Client

TCP Server



TCP socket: service setup

TCP Client

TCP Server

socket()

bind()

listen()

accept()

**blocked until
connection
from client**

TCP socket: service setup

TCP Client

socket()

TCP Server

socket()

bind()


listen()

accept()


**blocked until
connection
from client**

TCP socket: establish connection

TCP Client




TCP Server




**blocked until
connection
from client**

TCP socket: send and receive data

TCP Client




TCP Server



**blocked until
connection
from client**


TCP socket: send and receive data

TCP Client



data (request)


TCP Server



**blocked until
connection
from client**


TCP socket: send and receive data

TCP Client



data (request)

TCP Server




**blocked until
connection
from client**


process request

TCP socket: send and receive data

TCP Client



TCP Server



**blocked until
connection
from client**



process request

data (request)


data (reply)

TCP socket: close connection

TCP Client



TCP Server



**blocked until
connection
from client**



process request

data (request)




data (reply)

TCP socket: close connection

TCP Client



TCP Server



data (request)

data (reply)

blocked until
connection
from client

process request

Socket programming API: syscalls

- **int socket(int domain, int type, int protocol);**
 - Create a socket
 - returns the socket descriptor or -1(failure). Also sets errno upon failure
- **domain:** protocol family
 - **PF_INET** for IPv4, **PF_INET6** for IPv6, **PF_UNIX** or **PF_LOCAL** for Unix socket, **PF_ROUTE** for routing
- **type:** communication style
 - **SOCK_STREAM** for TCP (with **PF_INET**)
 - **SOCK_DGRAM** for UDP (with **PF_INET**)
- **protocol:** protocol within family, which is typically set to 0

Socket programming API: syscalls

- `int bind(int sockfd, struct sockaddr* myaddr, int addrlen);`
 - Bind a socket to a local IP address and port number
 - returns 0 on success, -1 and sets errno on failure
 - **sockfd**: socket file descriptor returned by `socket()`, `int` type
 - **myaddr**: includes IP address and port number
 - **NOTE**: `sockaddr` and `sockaddr_in` are of same size, use `sockaddr_in` and convert it to `socketaddr`
 - **sin_family**: protocol family, e.g. `AF_INET`
 - **sin_port**: port number assigned by caller
 - **sin_addr**: IP address
 - **sin_zero**: used for keeping same size as `sockaddr`
 - **addrlen**: `sizeof(struct sockaddr_in)`

```
struct sockaddr {
    short sa_family;
    char sa_data[14];
};

struct sockaddr_in {
    short sin_family;
    ushort sin_port;
    struct in_addr sin_addr;
    unsigned char sin_zero[8];
};
```

a pointer to a struct `sockaddr_in` can be cast to a pointer to a struct `sockaddr` and vice-versa

What's the difference between `PF_INET` and `AF_INET`???

Socket programming API: essential structs

- sockaddr — socket address info
- sockaddr_in — yet another struct for the ‘internet’

```
struct sockaddr {
    unsigned short sa_family; // addr family, AF_XXX
    char sa_data[14]; // 14 bytes of proto addr
};
struct sockaddr_in { // used for IPv4 only
    short sin_family; // addr family, AF_INET
    unsigned short sin_port; // port number
    struct in_addr sin_addr; // internet address
    unsigned char sin_zero[8]; // zeros, same size as sockaddr
};
struct in_addr { // used for IPv4 only
    uint32_t sin_port; // 32-bit IPv4 address
};
```

Socket programming API: syscalls

- **int listen(int sockfd, int backlog);**
 - Put socket into passive state (wait for connections rather than initiating a connection)
 - returns 0 on success, -1 and sets errno on failure
 - **sockfd**: socket file descriptor returned by socket()
 - **backlog**: the maximum number of connections this program can serve simultaneously

Socket programming API: syscalls

- **int accept(int sockfd, struct sockaddr* client_addr, int* addrlen);**
 - Accept a new connection
 - Return client's socket file descriptor or -1. Also sets errno on failure
 - **sockfd**: socket file descriptor for server, returned by socket()
 - **client_addr**: IP address and port number of a client (returned from call)
 - **addrlen**: length of address structure = pointer to **int** set to **sizeof(struct sockaddr_in)**
 - **NOTE: client_addr and addrlen are result arguments**
 - i.e. The program passes empty client_addr and addrlen into the function, and the kernel will fill in these arguments with client's information (**why do we need them?**)

More Information about Accept()

- A new socket is cloned from the listening socket
- If there are no incoming connection to accept
 - **Non-Blocking mode:** accept() returns -1 and throw away the new socket
 - **Blocking mode (default):** accept operation was added to the wait queue

Socket programming API: syscalls

- **int connect (int sockfd, struct sockaddr* server_addr, int addrlen);**
 - Connector to another socket (server)
 - Return 0 on success, -1 and sets errno on failure
 - **sockfd**: socket file descriptor (returned from socket)
 - **server_addr**: IP address and port number of the server
 - server's IP address and port number should be known in advance
 - **addrlen**: sizeof(struct sockaddr_in)

Socket programming API: syscalls

- **int write(int sockfd, char* buf, size_t nbytes);**
 - Write data to a TCP stream
 - Return the number of sent bytes or -1 on failures
 - **sockfd**: socket file descriptor from socket ()
 - **buf**: data buffer
 - **nbytes**: the number of bytes that caller wants to send

Socket programming API: syscalls

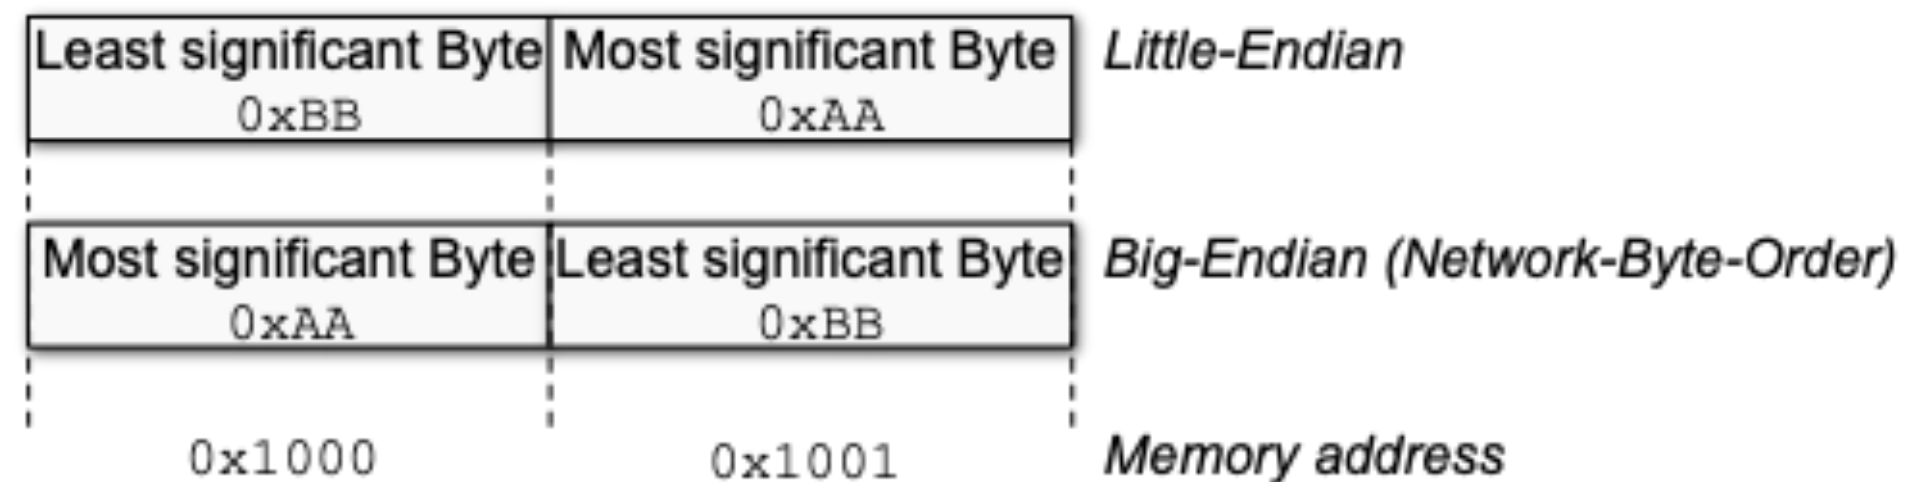
- **int read(int sockfd, char* buf, size_t nbytes);**
 - Read data from TCP stream
 - Return the number of bytes read or -1 on failures
 - Return 0 if socket is closed
 - **sockfd**: socket file descriptor returned from socket ()
 - **buf**: data buffer
 - **nbytes**: the number of bytes that caller can read (usually set as buffer size)

Socket programming API: syscalls

- **int close(int sockfd);**
 - close a socket
 - return 0 on success, or -1 on failure
 - After close, sockfd is no longer valid

Caveat: byte ordering matters

- Little Endian: least significant byte of word is stored in the lowest address
- Big Endian: most significant byte of word is stored in the lowest address
- Hosts may use different orderings, so we need byte ordering conversion
- **Network Byte Order = Big Endian**



Caveat: byte ordering matters

- Byte ordering functions: used for converting byte ordering

- Example: `int m, n;`
`short int s, t;`

<code>m = ntohl (n)</code>	<code>net-to-host long (32-bit) translation</code>
<code>s = ntohs (t)</code>	<code>net-to-host short (16-bit) translation</code>
<code>n = htonl (m)</code>	<code>host-to-net long (32-bit) translation</code>
<code>t = htons (s)</code>	<code>host-to-net short (16-bit) translation</code>

- Rule: for every int or short int
- Call `htonl()` or `htons()` before sending data
- Call `ntohl()` or `ntohs()` before reading received data

Address util functions

- All binary values are network byte ordered
- **struct hostent* gethostbyname (const char* hostname);**
 - Translate host name (e.g. “localhost”) to IP address (with DNS working)
- **struct hostent* gethostbyaddr (const char* addr, size_t len, int family);**
 - Translate IP address to host name
- **char* inet_ntoa (struct in_addr inaddr);**
 - Translate IP address to ASCII dotted-decimal notation (e.g. “192.168.0.1”)
- **int gethostname (char* name, size_namelen);**
 - Read local host’s name

FYI: struct hostent

<code>char *h_name</code>	The real canonical host name.
<code>char **h_aliases</code>	A list of aliases that can be accessed with arrays—the last element is <code>NULL</code>
<code>int h_addrtype</code>	The result's address type, which really should be <code>AF_INET</code> for our purposes.
<code>int length</code>	The length of the addresses in bytes, which is 4 for IP (version 4) addresses.
<code>char **h_addr_list</code>	A list of IP addresses for this host. Although this is a <code>char**</code> , it's really an array of <code>struct in_addr</code> 's in disguise. The last element is <code>NULL</code> .
<code>h_addr</code>	A commonly defined alias for <code>h_addr_list[0]</code> . If you just want any old IP address for this host (they can have more than one) just use this field.

Address util functions (cont'd)

- `in_addr_t inet_addr (const char* strptr);`

- Translate dotted-decimal notation to IP address (network byte order)

```
struct sockaddr_in ina;  
ina.sin_addr.s_addr = inet_addr("10.12.110.57");
```

- `int inet_aton (const char* strptr, struct in_addr *inaddr);`

- Translate dotted-decimal notation to IP address

```
struct sockaddr_in my_addr;  
my_addr.sin_family = AF_INET;           // host byte order  
my_addr.sin_port = htons(MYPORT);       // short, network byte order  
inet_aton("10.12.110.57", &(my_addr.sin_addr));  
memset(&(my_addr.sin_zero), '\0', 8); // zero the rest of the struct
```

Summary: what we have learned today

- What is the model for network programming?
 - **Client-Server model**
- Where are we programming?
 - **TCP and UDP in a nutshell**
- Which APIs can we use? How to use them?
 - **Socket programming**

Further Reading

- Stevens, W. Richard, Bill Fenner, and Andrew M. Rudoff. *UNIX Network Programming: The Sockets Networking API*. Vol. 1. Addison-Wesley Professional, 2004.
- Beej's Guide to Network Programming (<http://beej.us/guide/bgnet>)
- Socket Programming from Dartmouth,
<http://www.cs.dartmouth.edu/~campbell/cs60/socketprogramming.html>
- C/C++ reference: <http://en.cppreference.com>

Q&A

See you next time!

- TA: Qianru Li
- Office hour: 12:30-2:30pm, Monday

