

Project 2 overview

The best way to approach this project is in incremental steps. Do not try to implement all of the functionality at once.

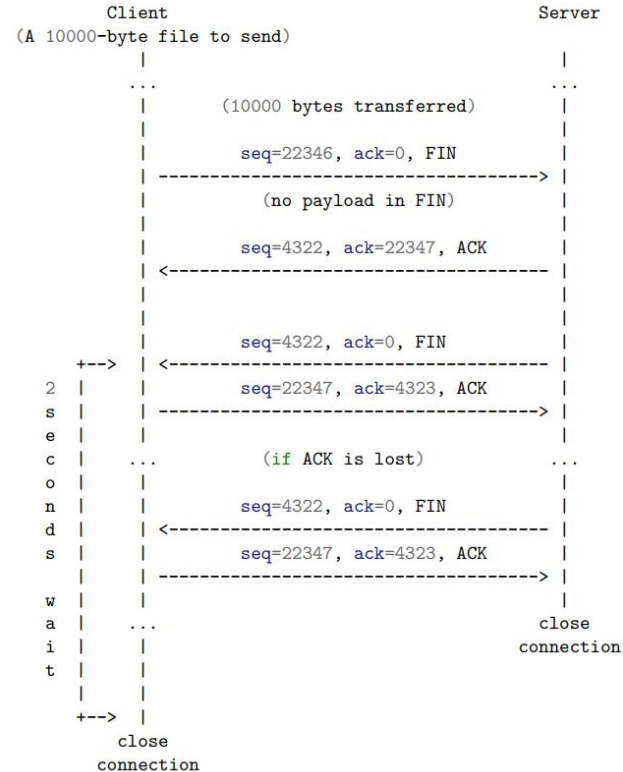
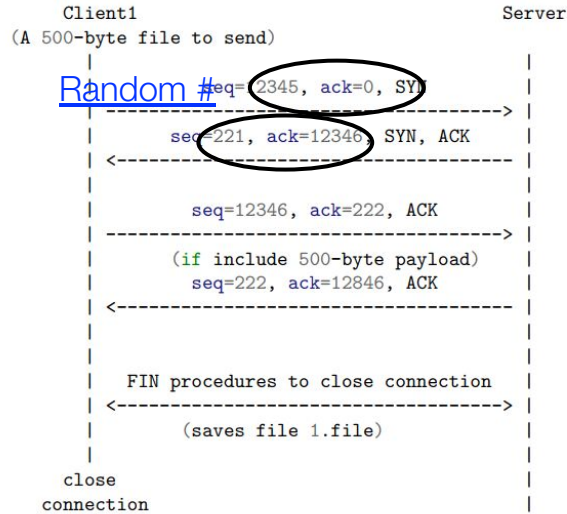
- First, assume there is no packet loss, implement the header fields and connection control functions (initialization with 3-way handshake and termination). Just have the client initiate the connection with 3-way handshake, send a small file (200 Bytes) as a packet, and the server respond with an ACK, and then the server use FIN procedure to close the connection.
- Second, introduce a large file transmission and pipe-lining. This means you must divide the file into multiple packets and transmit the packets based on the specified window size.
- Third, introduce packet loss. Now you have to add a timer for last sent packet (Go-Back-N) or several timers for each unacked packets (Selective repeat). If a timer times out, the corresponding (lost) packet should be retransmitted for the successful file transmission.

Stage 0: Small file transmission

- Small file transmission
 - A client initiate file transmission
 - A server accept connection requests, receive the file and save it with x.file
 - X indicated the counter of connection (starts with 1)
- Test
 - `./server 5000`
 - `./client localhost 5000 testfile`
 - In the server folder, check whether 1.file is saved and compare two files with diff command.

Stage 1: connection management

- Connection management



Stage 1: connection management

- Packet header struct (12 bytes)
 - Needed fields: a Sequence Number field, an Acknowledgment Number field, and ACK, SYN , and FIN flags.
 - Example :
 - uint16_t to represent each field.
 - In total, 5*2 bytes are used. Then pad 2 byte of zeros.
 - Functions: printPacket(), htonHeader(), ntohHeader().

Stage 1: connection management

- Packet header struct (12 bytes)
- Example to construct a SYN packet
 - Header h1, then memset the struct
 - Set sequence number fields of h1 with a random number
 - Set SYN flag
 - Print header “SEND 12345 0 SYN”
- Example to parse a packet
 - Print header “RECV 4321 12346 SYN ACK”

Stage 1: connection management

- Client side logic
 - Send a packet with SYN to initiate the connection.
 - After receive packet with ACK, start send packets with data.
 - After transmitting the entire file, send FIN packet and wait for ACK.
 - After receive server FIN, send ACK and wait for 2 seconds to close the connection.

Note: always need to print out the header

Stage 1: connection management

- Server side logic
 - If a SYN packet, reply with packet with SYN flag and ACK flag, set ACK number field and sequence number field
 - If a data packet, write data field to file
 - If a FIN packet, reply with packet with ACK flag. Then send a packet with FIN flag. After receive ACK from client, close the connection.

Note: always need to print out the header

Stage 2: large file transmission and pipelining

- Pipelining
 - For client side, send 10 packets at the same time.
 - For every received ACK, send a new packet out. Keep the window at 10.
 - For server side, no much difference
- Large file transmission
 - Pay attention to sequence number (max = 25600)

Stage 3: reliable data transfer with packet loss

- Go-back-N is recommended
- For client side
 - Keep a timer, restart the timer for every sent packet.
 - If timeout, resend all packets in the window.
- For server side
 - Keep expected sequence number
 - Every time a data packet is received, check whether the sequence number is expected. If expected, write data field, otherwise drop it.

Sample output during packet loss

At sender side:

SEND 100 0

RECV 123 200 ACK

SEND 200 0

SEND 300 0

SEND 400 0

TIMEOUT 200

RESEND 200 0

RESEND 300 0

RESEND 400 0

Timer implementation

- Logic of the timer
 - `setTimer(time period)`: Record the time out time point
 - `Bool <- isTimeout()`: Compare the current time with the time out time point. If smaller, return false; if equal or larger, return true
- How to get current time?
 - C++11 or higher
 - `auto now = std::chrono::high_resolution_clock::now();`
 - C
 - `gettimeofday(&timeval, NULL)`

Timer implementation

- How to compare time points
 - C++ 11 or higher: (**std::chrono::time_point** type)
 - if (now > timeout)
 - C: (**timeval** type)
 - double start = (double) s.tv_sec + (double) s.tv_usec/1000000;
 - double end = (double) e.tv_sec + (double) e.tv_usec/1000000;
 - If ((end - start) < 0.0)
- How to modify a time point
 - C++ 11 or higher: check +=, +, -, -= function of std::chrono::time_point
 - C: directly +=, +, -, -= on the timeval.tv_sec and timeval.tv_usec

Timer implementation

- You can use many other types from C or C++
 - E.g., you can simply use a large unsigned int `uint32_t` to keep POSIX timestamp to represent time
 - E.g., you can define your own time class/struct
- Tips
 - Time granularity is important (e.g., you shouldn't use second granularity to measure timeout)
 - Keep checking time out in your loop

Example of check timer to transmit SYN

Pseudo code:

```
While (1) {
    sendSYN();
    setTimer(); // set a 0.5 second timer
    if (sent) {printHeader(resend=true)};
    else {printHeader(resend=false)};
    sent=true;
    while (1) {
        // receive ack packet and break the loop if received.
        // if timer expires, break the loop.
        if (isTimeout()) break;
    }
    // If ack is received, break the loop.
}
```

Hint: how to make socket non-blocking

We set a flag on a socket which marks that socket as non-blocking. This means that, when performing calls on that socket, if the call cannot complete, then instead it will fail with an error like `EWOULDBLOCK` or `EAGAIN`.

- `fcntl()`: change the mode of file descriptor
 - Example: `fcntl(sockfd, F_SETFL, O_NONBLOCK);`
 - Note: `#include <fcntl.h>`

Tc command

Example of set loss and check it:

→ ~ sudo tc qdisc add dev lo root netem loss 10%

→ ~ sudo tc qdisc show dev lo

qdisc netem 8001: root refcnt 2 limit 1000 loss 10%

→ ~ ping localhost

PING localhost (127.0.0.1) 56(84) bytes of data.

64 bytes from localhost (127.0.0.1): icmp_seq=1 ttl=64 time=0.104 ms

64 bytes from localhost (127.0.0.1): icmp_seq=2 ttl=64 time=0.025 ms

..

64 bytes from localhost (127.0.0.1): icmp_seq=11 ttl=64 time=0.037 ms

64 bytes from localhost (127.0.0.1): icmp_seq=13 ttl=64 time=0.023 ms

...

^C

--- localhost ping statistics ---

22 packets transmitted, 19 received, 13% packet loss, time 20997ms

rtt min/avg/max/mdev = 0.018/0.031/0.104/0.020 ms